
Homework

Stochastic Models and Adaptive Algorithms

Réda Vince

Z697LX

January 4, 2021

Contents

1	Linear regression	2
1.1	Function approximation with least squares	2
1.2	Approximating auto-regressive series	4
2	Kernel methods	5
3	Reinforcement learning	6

1 Linear regression

1.1 Function approximation with least squares

Ordinary least-squares

The best linear approximation of a function can be calculated using least-squares.

At first, a noisy sample of (x, y) pairs is generated such that $y_i = c x_i \sin(c x_i) + \epsilon_i$, where $\epsilon_i \sim \mathcal{N}(0, 1)$.

Let $[\Phi]_{ij} = f_j(x_i)$ be the transformed input vector and \mathbf{y} the output, where f_j is a basis function. From this the $\Phi(x)$ matrix can be generated after selecting a suitable f . A number of these were tried, and the best one for the problem seemed to be the polynomial one, that is $f_i(x) = x^{i-1}$. As a parameter, $d = 10$ was used.

Now we have to find the optimal $\hat{\theta}$ parameter vector, for which $\Phi \theta = \mathbf{y} = [y_1 \dots y_n]^T$. This is done like so: $\theta^* \approx \hat{\theta} = \Phi^+ \mathbf{y}$. The Φ matrix of the sampled inputs and of the LS estimate are the same, so the function is evaluated at the same x values.

A computationally cheaper method for the pseudoinverse is QR decomposition. For this, the "economic" mode of scipy's `qr` function is used. Then the pseudoinverse is $\Phi^+ = \mathbf{R}^{-1} \mathbf{Q}^T$. Because the pseudoinverse of a matrix is unique, this method gives the same result as the previous one.

The results are shown in figure 1a.

Recursive least-squares

Next, more of the above described (x, y) pairs is sampled and $\hat{\theta}_n$ calculated periodically using recursive least-squares. In the code, all of the samples are measured beforehand for simplicity.

The equation for $\hat{\theta}$ can be reformulated in the following way:

$$\hat{\theta} = \underbrace{\left[\sum_{i=1}^n \varphi \varphi^T \right]^{-1}}_{\Psi_n} \underbrace{\sum_{i=1}^n y_i \varphi_i}_{z_n}. \quad (1)$$

Now we have an update rule for both $\Psi_{n+1} = (\Psi + \varphi_{n+1} \varphi_{n+1}^T)^{-1}$, and $z_{n+1} = z_n + \varphi_{n+1} y_{n+1}$. Both Ψ_0 and z_0 are set to zero. Also, in the code $\hat{\theta}_n$ is calculated only when plotted for speed.


The resulting plots are shown in figure 1c., taken at $n \in [25, 50, 75, 100]$.

Least-norm problem


Now let's make $d > n$, specifically $n = 100$ and $d = 200$, which makes Φ fat. We make the assumption that Φ is still full-rank. The solution is the same, except we are going to use singular value decomposition (SVD) for the pseudoinverse of Φ .

The SVD is calculated like this: $\Phi_{d \times n} = \mathbf{U}_{d \times d} \mathbf{\Sigma}_{d \times n} \mathbf{V}_{n \times n}^T$. Let's denote the matrix of column vectors of the normalized eigen-vectors of matrix Φ by $\text{eig}(\Phi)$. Let's denote the eigenvalues by $\text{eigval}(\Phi)$. Then, $\mathbf{U} = \text{eig}(\Phi \Phi^T)$, $\mathbf{V} = \text{eig}(\Phi^T \Phi)$ and $\mathbf{\Sigma} = \text{diag}(\text{eigval}(\Phi \Phi^T))_{d \times n}$. Then, the pseudoinverse is $\Phi^+ = \mathbf{V} \mathbf{\Sigma}^+ \mathbf{U}^T$. For $\mathbf{\Sigma}^+$, we take the inverse of the non-zero elements of $\mathbf{\Sigma}$, and add zeros such that it has the shape of $d \times n$.


The resulting plots are shown in figure 1b.



(a) Least-squares estimate for thin Φ



(b) Least-norm estimate for fat Φ



(c) Recursive least-squares

Figure 1: Experiments with ordinary least-squares

1.2 Approximating auto-regressive series

A recursive time-series is generated from the give equation: $y_t = a y_{t-1} + b y_{t-2} + \epsilon_t$. Let's calculate the least-squares estimate using $\varphi_t = [y_{t-1}, y_{t-2}]$, $\Phi = [\varphi_1 \dots \varphi_n]$. Then $\hat{\theta} = \Phi^+ y$.

The time plots can be seen on figure 2a.

Let's calculate the inverse of the covariance matrix:s $\Gamma_n = \frac{1}{n} \Phi^T \Phi$. Now define $\Delta\theta := (\theta - \hat{\theta}_n)$. The confidence ellipsoid is given by

$$\Delta\theta^T \Gamma_n \Delta\theta \leq \frac{q \hat{\sigma}_n^2}{n}, \quad (2)$$

where q is calculated from the inverse of the cumulative χ^2 distribution function given the p probabilities ($q = F(p)_{\chi^2(d)}^{-1}$). In this problem, $d = 2$. Eq. 2 means that with probability p , the optimal θ^* is at most $\Delta\theta$ distance from $\hat{\theta}$.

Now we assume that $\hat{\sigma}_n = 1$, and let $\Gamma_{ij}/n = [\Gamma_n]_{ij}$. Then we have an equation that outputs an ellipse for a p probability. Written out:

$$\Delta\theta_1^2 \Gamma_{11} + 2 \Delta\theta_1 \Delta\theta_2 \Gamma_{12} + \Delta\theta_2^2 \Gamma_{22} = q. \quad (3)$$

For the plotting, this equation is transformed so that the axis distances and the rotation angle is known with the function `ellipse_transform[1]`.

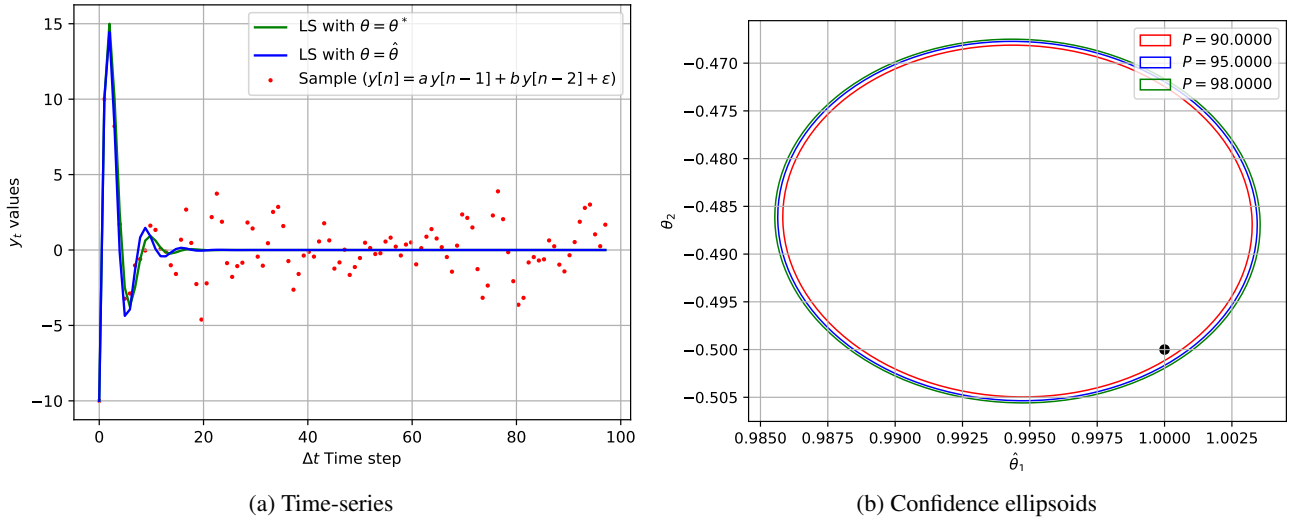


Figure 2: Experiments with auto-regressive series

2 Kernel methods

3 Reinforcement learning

References

- [1] Ellipse equation: https://www.maa.org/external_archive/joma/Volume8/Kalman/General.html#:~:text=Applying%20the%20methods%20of%20%E2%8C%99%BD,rotated%20through%20an%20angle%20%E2%8C%B1%20.&text=which%20is%20in%20the%20form,with%20A%20and%20C%20positive.