



Stochastic Models & Adaptive Algorithms

Balázs Csanád Csáji

SZTAKI: Institute for Computer Science and Control
13-17 Kende utca, Budapest, Hungary, 1111

Ph.D. Course, BME & ELTE, Semester: 2020/21/1

Main Parts of the Course

I. STATISTICAL LEARNING (40 %)

From linear regression to kernel methods. Performance guarantees. This part mainly addresses **regression** and **classification** problems.

II. REINFORCEMENT LEARNING (40 %)

Learning via interactions with an uncertain, dynamic environment. **Markov** decision processes and model-free **online learning** methods.

III. STOCHASTIC APPROXIMATION (20 %)

A general theory of recursive methods working with uncertain data. E.g., stochastic **optimization**, acceleration techniques, **consistency**.

Assumed Mathematical Background

LINEAR ALGEBRA

Positive (semi-) definite matrices, QR decomposition, eigenvalues, eigenvectors, pseudoinverse, matrix norms, range and null space.

MATHEMATICAL ANALYSIS

Convex functions, gradient vector, Hessian matrix, Lagrange multipliers, Hilbert spaces, Riesz representation theorem.

PROBABILITY & STATISTICS

Central moments, convergence concepts (a.s., in distribution, etc.) bias (of an estimator), Fisher information, statistical efficiency.

*"One must learn by doing the thing;
for though you think you know it,
you have no certainty, until you try."*

(Sophocles)

NUMERICAL EXPERIMENTS

PART I: LINEAR REGRESSION

Numerical Exercise I.1

- Generate a **sample** of size n , $\{(x_i, y_i)\}_{i=1}^n$, where $\{x_i\}$ are i.i.d. random variables having a uniform distribution on $[0, 1]$; let $y_i = f(x_i) + \varepsilon_i$, where $f(x) = x \sin(cx)$, e.g., with $c = 16$, and the **noises**, $\{\varepsilon_i\}$, are zero-mean i.i.d. Gaussian random variables.
- Construct **linear regression** estimates with d **basis** functions using the **least-squares** criterion. Experiment with different types of basis functions, such as $f_i(x) = x^{i-1}$ and $f_i(x) = \exp(-\|x - \mu_i\|^2 / \sigma_i^2)$; e.g., one can use $\mu_i = (i - 1)/(d - 1)$ and a fixed constant σ^2 for σ_i^2 .
- Calculate the LS estimate with **QR** decomposition or with **SVD**.
- Apply the **recursive** formulation of LS for increasing sample sizes.
- Also make experiments for the **least-norm** problem, i.e., if $d > n$.
- **Plot** the observations, the original function, and the resulting regression functions for various sample sizes and basis choices.

Numerical Exercise 1.2

- Consider the scalar **AR** (auto-regressive) **time-series** model,

$$y_t = a \cdot y_{t-1} + b \cdot y_{t-2} + \varepsilon_t,$$

where $\{\varepsilon_t\}$ are i.i.d. standard normal; a and b are parameters (constants) with $b < 1 + a$, $b < 1 - a$ and $b > -1$ (for stability).

- This system can be written in a **linear regression** form as

$$y_t = \varphi_t^T \theta^* + \varepsilon_t,$$

where $\varphi_t = (y_{t-1}, y_{t-2})^T$, the rows of Φ , and $\theta^* = (a, b)^T$.

- The error of the LS estimates is **asymptotically Gaussian**, that is

$$\sqrt{n}(\hat{\theta}_n - \theta^*) \xrightarrow{d} N(0, \Gamma) \quad \text{as } n \rightarrow \infty,$$

where $\hat{\theta}_n$ is the LS estimate for sample size n , “ \xrightarrow{d} ” denotes convergence in distribution and $N(0, \Gamma)$ is the normal distribution.

Numerical Exercise 1.2

- The asymptotic **covariance matrix** Γ is $\sigma_0^2 \mathbb{E}[\varphi_0 \varphi_0^T]$, where σ_0^2 is the variance of the noises, $\{\varepsilon_t\}$ (i.e., one for standard normal), and φ_0 has the stationary distribution of the regressors, $\{\varphi_t\}$.
- Assuming $\sigma_0^2 = 1$, this can be estimated by (cf. ergodic theorem)

$$\sigma_0^2 \mathbb{E}[\varphi_0 \varphi_0^T] \approx \Gamma_n = \frac{1}{n} \sum_{t=1}^n \varphi_t \varphi_t^T.$$

- Choose an appropriate θ^* , some initial conditions (y_0, y_{-1}) and simulate a **trajectory** of size $n = 100$. Then, **plot** the LS estimate with its **confidence ellipsoids** for probabilities $p = 0.9, 0.95, 0.98$.
- Of course, one should build the ellipsoids by using the fact that

$$n(\hat{\theta}_n - \theta^*)^T \Gamma_n^{-1} (\hat{\theta}_n - \theta^*) \sim \chi^2(\dim(\theta^*)),$$

where $\chi^2(d)$ is the χ^2 distribution with d degrees of freedom.

Requirements

- Perform the experiments using **Matlab** or **Python**. Adequately structure and **comment** your codes, to make them easily readable.
- Use **matrix** manipulating functions for QR decomposition, SVD, inverse, etc.; but do **not** use existing functions for pseudoinverse or model fitting, such as the backslash operator of Matlab for LS.
- Write a short, 1-3 pages, **report** with figures which describes and illustrates your experiments. Make the report in \LaTeX or Jupyter.
- Send the codes and the report in **email** to csaji@sztaki.hu.
- Hint 1: if a square matrix A is badly conditioned, a simple way to decrease its condition number is to replace it with $A + \lambda I$, where I is the identity matrix and λ is a tiny scalar (cf. ridge regression).
- Hint 2: one can look up the inverse CDF values of $\chi^2(d)$ from a table, but there are also functions in Matlab and Python to do it.

*"I hear, I forget;
I see, I remember;
I do, I understand."*
(Chinese proverb)

NUMERICAL EXPERIMENTS

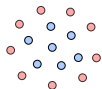
PART II: KERNEL METHODS

Numerical Exercise II.1

- Consider a **binary classification** problem on the plane, \mathbb{R}^2 , where the data are generated as follows: for each observation, for both classes, there is a 0.5 probability of getting the observation from that class; and the classes have multivariate normal distributions.
- The **means** of the distributions should be different (mandatory), and their **covariance** matrices could also be different (optional).
- Generate datasets and implement the following **linear classifiers**: **soft margin SVC**, **LS-SVC** and **nearest centroid classifier** (NCC).
- For the first two cases, use and solve the **primal** form of the corresponding optimization problems (do not use SVM libraries).
- **Plot** the datasets and the obtained **classification boundaries** (for SVC, LS-SVC and NCC) for various regularization parameters (λ) and various datasets (generated by different means/covariances).

Numerical Exercise II.2

- Consider a **binary classification** problem on the plane, \mathbb{R}^2 , where one of the classes mostly concentrates in a compact area, while the other class surrounds the previous class (with potential overlaps).



- First, design an algorithm to generate such classification datasets.
- Implement the **nonlinear** (kernelized) **soft margin SVC** method using the **Gaussian kernel**, namely, $k(x, y) = \exp(-\|x - y\|^2 / \sigma^2)$.
- For this, the **Wolfe dual** of soft margin SVC should be solved by a convex optimization tool (do not directly use SVM libraries).
- **Plot** the datasets and the obtained SVC **classification boundaries** for various regularization parameters (λ) and kernel widths (σ).

Requirements

- Perform the experiments using **Matlab** or **Python**. Adequately structure and **comment** your codes, to make them easily readable.
- Use a **convex optimization** library, e.g., **cvx** in Matlab or **cvxpy** in Python, to implement the corresponding primal or dual problems.
- Important: do **not** use existing methods dedicated to SVMs, such as the ones provided by the LibSVM library in Matlab and Python.
- The aim is to implement SVC based on mathematical optimization tools and not to simply use an existing SVM library for that.
- Also, do not copy SVC from a source code found on the internet.
- Write a short, 1-3 pages, **report** with figures which describes and illustrates your experiments. Preferably, make the report in \LaTeX . (you can also make the documentation in Jupyter notebook).
- Send the codes and the report in **email** to csaji@sztaki.hu.

*"Being a student is easy.
Learning requires actual work."*
(William Crawford)

NUMERICAL EXPERIMENTS

PART III: REINFORCEMENT LEARNING

Numerical Exercise III.1

- Consider a small-scale **Markov Decision Process** (MDP), such as the gambler's problem, cliff walking or even a randomly generated MDP (e.g., see Examples 4.3 and 6.6 of Sutton and Barto's book: "Reinforcement Learning: An Introduction", 2nd. edition, 2018).
- Implement the following **model-based** MDP solution algorithms: **value iteration** (VI), **policy iteration** (PI), **linear programming** (LP).
- Compute the **optimal** value function of the selected MDP with LP.
- Then, plot its **distance** (using Euclidean distance, i.e., $\| \cdot \|_2$) from the value functions calculated by the VI and PI algorithms.
- The comparison should be **per iteration** (hence, the x-axis of the figure should show the iteration number, while the y-axis should present the distances of VI and PI from the optimal solution).

Numerical Exercise III.2

- Consider the same small-scale MDP as in the first experiments.
- Implement the **model-free** Q-learning using a totally (uniformly) **random**, an **ε -greedy** and a **Boltzmann** (semi-greedy) base policy.
- Thus, you should simulate the environment using a base policy and iteratively learn the optimal action-value function by applying Watkins' Q-learning on the obtained state-action-cost trajectories.
- Hints: ε could be 0.1 or 0.05, while τ could be 0.1, 1, and 10.
- Let V^* be the optimal value function (computed by LP in the first experiments) and let $V_t(x) \doteq \min_a Q_t(x, a)$ be the value function induced by a Q-learning algorithm (for iteration t). Plot the **distances** ($\|\cdot\|_2$) of V^* and V_t for the three Q-learning approaches.
- Also plot (on a different figure) the gathered **costs/rewards** during the learning process for the case of all three base policies above.

Requirements

- Perform the experiments using **Matlab** or **Python**. Adequately structure and **comment** your codes, to make them easily readable.
- Use a **convex optimization** library, e.g., **cvx** in Matlab or **cvxpy** in Python, to implement the resulting linear programming problem.
- Important: it is **allowed** to use existing **RL environments** for the MDPs, but not for the solution algorithms. Do **not** use existing methods dedicated to RL, **implement** the **RL algorithms** yourself.
- The aim is to understand the fundamental RL approaches by implementing them, and not simply to apply RL codes of others.
- Write a short, 1-3 pages, **report** with figures which describes and illustrates your experiments. Preferably, make the report in \LaTeX . (you can also make the documentation in Jupyter notebook).
- Send the codes and the report in **email** to csaji@sztaki.hu.

"The true method of knowledge is experiment."
(William Blake)

NUMERICAL EXPERIMENTS

PART IV: STOCHASTIC APPROXIMATION

Numerical Exercise IV.1

- Consider the **Q-learning** algorithm implemented in Num. Ex. III.2 (any of the three base policies could be used for this exercise).
- Study the extension of Q-learning with **momentum acceleration** and **Polyak averaging**. If γ_n and β_n are the stepsizes of Q-learning and the momentum term, resp., consider using $\gamma_n = c_1/n(x, a)$ and $\beta_n = c_2/n(x, a)$, where $n(x, a)$ is the number of visits to (x, a) .
- Important: when calculating the momentum terms and the Polyak averages of the estimates, the $Q_n(x, a)$ values should be treated as **different** stochastic approximation processes **for each (x, a)** .
- Namely, only those $Q_n(x, a)$ values should be taken into account for the momentum or averages, where that (x, a) pair was visited.
- **Compare** the standard Q-learning with variants using momentum acceleration (study different c_i 's). For both approaches, also make the comparison if one applies Polyak averaging using a fix window.

Requirements

- Perform the experiments using **Matlab** or **Python**. Adequately structure and **comment** your codes, to make them easily readable.
- Use the MDP studied in Numerical Exercise III, therefore, it is **allowed** to use existing **RL environments** for the MDP model.
- Extend your implementation of Q-learning from Num. Ex. III.2. Do **not** use existing RL libraries, **implement** the methods yourself.
- Write a short, 1-3 pages, **report** with figures which describes and illustrates your experiments. Preferably, make the report in \LaTeX . (you can also make the documentation in Jupyter notebook).
- It should also contain **plots** of iterations vs distances from V^* , as in Num. Ex. III.2, for the algorithms (Q-learning, QL with momentum, and Polyak averages of both; preferably for various (c_1, c_2) pairs).
- Send the codes and the report in **email** to csaji@sztaki.hu.

INFO FOR THE EXAM

OVERVIEW OF THE KEY CONCEPTS

Question Groups for the Exam

- **Deterministic linear regression**: normal equations, orthogonal projections, recursive LS, least-norm problem, pseudoinverse, SVD
- **Stochastic linear regression**: Gauss-Markov theorem, ML vs LS, strong consistency, limiting distribution, PCA, ridge regression
- **Nonlinear optimization**: Lagrangian, weak and strong duality, Karush-Kuhn-Tucker conditions, Slater's condition, Wolfe duality
- **Statistical learning theory**: classification, true vs empirical risk, Bayes classifier, VC dim., ERM vs SRM, hard/soft margin SVM
- **Kernel methods**: Wolfe dual of SVMs, RKHS, representer theorem
- **Markov decision processes**: value functions, Bellman operators, value and policy iteration, linear programming, policy gradient
- **Reinforcement learning**: Monte Carlo, $TD(\lambda)$, SARSA, Q-learning
- **Stochastic approximation**: Robbins-Monro, Kiefer-Wolfowitz, SPSSA, momentum acceleration, Polyak averaging, Lyapunov functions

Thank you for your attention!

 www.sztaki.hu/~csaji

 csaji@sztaki.hu